



Tenderfoot

Behind the explainer — what is actually built, and where the team is

ADDENDUM FOR MANAGEMENT

16 SEPTEMBER 2026

COMPANION TO THE EXPLAINER OF AUGUST 2026

1 • WHAT THE EXPLAINER PROMOTES

The explainer sells the destination; this document says what is standing on the site today, and what the crew is doing next. It was written to show the **finished product** rather than the first release — its own closing page says so. Five weeks on, each promise against what exists:

PAGE	THE PROMISE	WHERE IT STANDS TODAY
03 What it does	It collects from every source you switch on, prepares the documents and facts, and you decide with a reason.	BUILT All three moves exist and have been used on real data.
04 The daily driver	One opportunity at a time, keyboard-only. Interested, Pass, Undo.	BUILT • USED IN ANGER 150 real opportunities were triaged in the app on 13 September.
05 Evidence	Scores that quote the sentence they came from.	PARKED, BY DESIGN V1 has no scores (page 2). The citation mechanism does exist for extracted facts: every deadline points to its passage.
06 The brief	What the work is, why it fits, what is missing.	NOT BUILT The record screen exists with its Extracted Fields, Documents and Timeline tabs; Brief is shown as parked.
07 The headline feature	Contracts expiring, months before an RFP.	DATA LOADED • SCREEN NOT BUILT All 204,920 Indiana contracts are in production; the radar screen is scheduled after the go / no-go.
08 Recall	Rejected is filed with its reason, never deleted.	PARTLY Every decision carries a reason in the reader's own words and can be undone. The drawer of machine-rejected items has nothing to hold: V1 rejects nothing on your behalf.
09 Configuration	Your firm is a row; sources are the control.	BUILT The firm profile is editable data; the Source Registry shows every source with its health, and each can be checked or run from the screen.

TWO THINGS TO SAY PLAINLY ABOUT THE EXPLAINER ITSELF

Its screenshots carry a placeholder wordmark and predate the current prototype, so it stays internal until rebuilt. And its closing “Where this stands” page describes the plan as it was in August; page 3 here replaces it.

What is actually running today

A deployed web application with a real database behind it, built since 3 August: **607 commits**, **over 1,100 automated tests**, and a CI gate that must pass before anything merges.

01 It collects — from three live sources

One solicitation seen on two sources becomes one record. Nothing is filtered out on the reader's behalf.

9,883

SAM.gov

Federal solicitations, pulled directly and free. In production.

4,377

HigherGov

Paid aggregator: IN, OH, MI, IL, KY at state *and* local level. 851 buying organisations; a complete 28-day window bought so coverage could be measured, not guessed.

204,920

Indiana contract register

Every contract the state holds, with value and end date. The data behind the expiration radar — loaded, waiting for its screen.

02 It prepares

When a reader opens a solicitation, its documents are fetched and read, and facts — the deadline first — are extracted with a confidence and the passage they came from. First live measurement: the deadline was **right every time it was found (100% precision)** and **found less often than we want (recall 12.5%, a lower bound not yet validated by hand-labelling)**.

03 You decide

The queue is real. One opportunity, keyboard-only, Interested or Pass with a required reason. The reasons are captured verbatim, and they are the raw material for the next design step (page 4).

And it administers itself: a Source Registry with a health check and a run control per source; a firm profile held as data, not code; a spend ledger for the one metered source.

WHAT IT DELIBERATELY DOES NOT DO

- **It does not rank, score or filter.** V1 returns everything from every switched-on source. Decided 11 August, and a reason rather than an omission: a system that returns everything cannot silently miss anything, and nobody knew what the sources actually publish until it ran. That measurement is now in hand.
- **It does not run unattended.** Every ingest is started by an operator. Scheduling is the first thing after the go / no-go.
- **It does not yet show the radar, the brief, organisations, or reports.** Those destinations exist in the navigation and say honestly that they are not built.

Where the team is

The short version. The plumbing is built and proven on real data. On 3 September the project turned from building screens to measuring the sources they depend on; that measurement is nearly complete. One decision sits with management, one gate with the team.

How we got here

August was construction — schema, a design system matched pixel-for-pixel to the prototype, federal ingestion, cross-source merging, document extraction, the triage screen — each slice merged with its tests. By month's end the app was deployed and taking real federal data.

“Once we have both our reliable upstreams defined, we can then flesh out the rest of the app with real data and analysis methods, instead of just hoping that the source meets our criteria.”

MATT • 3 SEPTEMBER • THE PIVOT

The pivot

Measured, the free state portals were thin — Indiana's publishes no archive, Kentucky's and Michigan's show only what is open today, Ohio's can only be read by hand. Rather than build screens over a weak foundation, the team built a **data floor** — seven pass/fail tests the data must meet before a go / no-go means anything — and evaluated a paid aggregator against a free answer key.

The finding

HigherGov, on a two-week trial, recalled nearly every state-agency notice with five weeks of bidding time still left when it first carried them, and holds an Indiana archive back to 2013 that the state itself does not publish. **That is a buy at \$500 a year.** One honest caveat: it was bought for city, county and school-district coverage, and that segment could not be measured — there is no free source to check it against.

The gate, and the cost discipline

The data floor now passes six of its seven tests; the seventh — where a listing defers to a document, we hold the document — can only close through use, because documents are fetched on demand to protect the metered budget. In practice, **the go / no-go is ready to be taken.** That budget, 10,000 records a month, is the project's one metered thing, and two rules bind every call: nothing is fetched without approval outside a small standing allowance, and rejection must be free — documents are bought only for what survives its listing. A key that leaked on 3 September was rotated within the hour, and every response is now scrubbed at the boundary.

THE FOUR NUMBERS TO QUOTE

98.7–100%

of state-agency notices recalled by HigherGov against a free answer key.

37–39 days

of bidding time still left, median, when HigherGov first carries a notice.

6 of 7

data-floor tests passing. The seventh closes only through use.

8,682/10,000

metered records spent by the trial's close; the rest deliberately unspent.

In order, each gated on the one before it

01 • NOW • FREE

Turn the 150 reasons into a vocabulary

The triage session produced 150 decisions (11 Interested, 139 not) in the reader's own words. The fixed reasons the queue will offer — “out of state”, “deadline passed” and their siblings — are derived from those, not invented. Needs Matt at the table.

02 • NOW • FREE

Two small rulings

Surfaced while building the last set; neither costs anything. One narrows a ruled design, the other repairs a measurement rule that no longer mirrors the one it was built to mirror.

03 • THE GATE

Go / no-go

With the floor at six of seven and the sources graded, the team can put the question: does the data justify the post-gate slices?

04 • ON GO

Live ingestion

Scheduled runs, and health alarms that fire when a source goes quiet. This is what turns an instrument into a daily tool.

05 • ON GO

The radar and reporting

The expiration radar over the 204,920 contracts already held; source yield and market sizing.

06 • LATER, DELIBERATELY

Everything parked with a reason

Scoring and evidence, the brief, organisations and vendors, the pipeline board, and sign-in — each parked with a reason on record, none forgotten.

DECISIONS MANAGEMENT CAN MAKE NOW

- **The HigherGov subscription.** The trial has ended. Renewal at \$500/yr is recommended on the evidence on page 3, with the sub-state caveat stated.
- **Whether the explainer goes outside the firm.** Not until it is rebuilt against the current prototype and its closing page is replaced with this one.

RISKS WORTH KNOWING

- Extraction recall is unvalidated.
- The sub-state coverage the aggregator was bought for is unmeasured.
- Nothing runs without a person starting it.
- The app is public behind platform deployment protection only, with no sign-in yet.

The explainer promised a tool that removes every reason the bid decision is slow. What exists today does the collecting and the preparing on real data, at a known cost, and has been used to decide. The rest is sequenced behind a gate that is ready to be taken.

Every figure here is traceable to the project's own records as of 16 September 2026: STATUS.md, the session log, the plan of action, and the fitness and platform-comparison documents under docs/.